

# QuestDB

## Community Meetup #7

2022-07-28

QuestDB Team



# Today's agenda

1 Update about replication

2 AMA

3 Other news





# Roadmap updates

Javier Ramirez, Developer Relations at QuestDB



@supercoco9 - @QuestDB



# QuestDB

## Recent developments

- Replication (WIP) ← **today's focus!**
- Partitions detach/attach workflow (soon)
- CSV imports for unordered data without size limitation (soon)



# What do we mean by replication

- Higher availability for reads, and higher throughput
  - Write to **one** QuestDB instance, propagate changes, and read from **many**
- Higher availability and higher throughput for both reads and writes (available only on QuestDB Cloud)
  - Write to **many** QuestDB instances, propagate changes, and read from **many**



# Current model for writing table data

- Data is received and kept in memory
- Once CommitLag is reached (or maxuncommitted rows, or table idle...), the TableWriter writes data into the table files
- If Out-of-order rows are present, the operation will be costly, affecting reads
- Depending on how costly the writing operation is, ingestion performance will also be affected
- Prone to Table Busy errors

# First Step: Detaching ingestion and table writes



Implementing WAL (Write Ahead Log) (Work In Progress) ⚠️

- Data (and any updates) is ingested to an append-only file (multiple WAL per table)
- If we take an empty DB and apply the changes in the WAL, we should get to the current dataset
- WAL file is structured by columns and divided into segments
- Segments have sequential IDs and are tracked by a TableSequencer, who coordinates and avoids conflicts
- Periodically, changes contained in the WAL are committed by the TableWriter



# WAL benefits

- Enables replication
- Decouples ingestion and writing, allowing for future optimizations
- WAL will remove ingestion slow-down in the case of Out-of-order data
- WAL will improve the TableBusy scenarios





## Second step: WAL replication and read replicas

(Expected in a few months from now. Depends on WAL)

- Each WAL segment will automatically be written into Amazon S3
- On S3 notification, the TableSequencer will fetch the pending segments in the right order and will use the TableWriter to apply changes locally
- The replicated tables will achieve eventual consistency
- The unit of replication will be the table
- Initially replication using only S3. If you want to contribute by making the TableSequencer interact with segments in other Object Storages (cloud or not), do get in touch



# AMA

Javier, Developer Advocate  
& QuestDB Team



# Questions from the community

- Why does it take a while to see data after ingesting?
- Where can I find the Grafana plugin for QuestDB?
- What timestamp format do I need to use with ILP?



# We're hiring!

- Senior Backend Engineer (Python), QuestDB Cloud
- Cloud Engineer, QuestDB Cloud

Apply at: [questdb.io/careers](https://questdb.io/careers)



# Thank you!



@QuestDb



QuestDB

Join Slack: [slack.questdb.io](https://slack.questdb.io)

Star us: [github.com/questdb/questdb](https://github.com/questdb/questdb)

# Clustering

## Diagram

